

Tutorial by Nooh Ayub

(SSD for Face Mask Detection Using Pretrained Model)

Objectives:

- Understand Pre-trained SSD model using different backbones (VGG16, Resnet34).
- train this pretrained model for your custom dataset (FaceMask).
- Evaluating the model on face mask test dataset and comparing the performance of different SSD backbones (VGG16, Resnet34)

Here's a step by step breakdown:

Step 1 — Create a directory named “my_ssd_implementation” and inside this directory create a virtual environment:

(a) create a conda environment with python 3.9:

```
conda create --name ssd_custom python=3.9|
```

(b) now activate this environment:

```
conda activate ssd_custom█
```

Step 2 — Now copy the contents of “project_ssd” inside the parent directory “my_ssd_implementation”:

Note: the project folder is present in the zip file

Step 3 — Next install Pytorch 1.13.0 with cuda 11.7 support from the official Pytorch website:

```
pip install torch==1.13.0+cu117 torchvision==0.14.0+cu117 torchaudio==0.13.0 --extra-index-url https://download.pytorch.org/whl/cu117
```

Step 4 —Install the dependencies from the requirements.txt

≡ requirements.txt.txt ×

≡ requirements.txt.txt

```
1 matplotlib==3.9.3
2 albumentations==1.0.3
3 torchmetrics==0.10.3
4 numpy==1.26.4
5 tqdm==4.67.1
6 opencv-python
```

Use the following command when installing the dependencies:

```
pip install -r requirements.txt
```

Step 5— copy the data folder from the zip file and paste it in this parent directory “my_ssd_implementation”:

(a) the dataset must be in PASCALVOC-2007 format.

(b) The dataset tree structure is as follows :

```
1. |— Test
2. |   |— Test
3. |       |— JPEGImages [470 entries exceeds filelimit, not opening dir]
4. |— Train
5. |   |— Train
6. |       |— JPEGImages [1888 entries exceeds filelimit, not opening dir]
7. |— Val
8. |   |— Val
9. |       |— JPEGImages [320 entries exceeds filelimit, not opening dir]
```

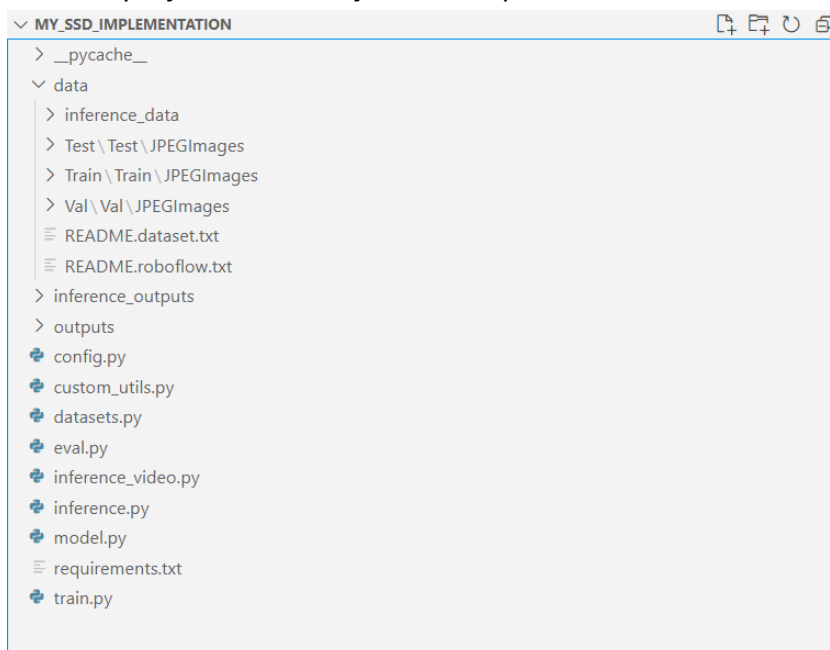
(c) The whole Parent directory structure should look like the following:

```

1. | data
2. | | inference_data
3. | | | image_1.jpg
4. | | | ...
5. | | | image_4.jpg
6. | | | video_1.mp4
7. | | Test
8. | | | Test
9. | | | | JPEGImages [470 entries exceeds filelimit, not opening dir]
10. | | Train
11. | | | Train
12. | | | | JPEGImages [1888 entries exceeds filelimit, not opening dir]
13. | | Val
14. | | | Val
15. | | | | JPEGImages [320 entries exceeds filelimit, not opening dir]
16. | inference_outputs
17. | | images [239 entries exceeds filelimit, not opening dir]
18. | | videos
19. | | | video_1.mp4
20. | notebooks
21. | | visualizations_data.ipynb
22. | outputs
23. | | best_model.pth
24. | | last_model.pth
25. | | map.png
26. | | train_loss.png
27. | config.py
28. | custom_utils.py
29. | datasets.py
30. | eval.py
31. | inference.py
32. | inference_video.py
33. | model.py
34. | train.py

```

(d) Final display structure of your workspace:



The screenshot shows a file explorer window titled 'MY_SSD_IMPLEMENTATION'. The directory structure is as follows:

- __pycache__
- data
 - inference_data
 - Test \ Test \ JPEGImages
 - Train \ Train \ JPEGImages
 - Val \ Val \ JPEGImages
 - README.dataset.txt
 - README.roboflow.txt
- inference_outputs
- outputs
- config.py
- custom_utils.py
- datasets.py
- eval.py
- inference_video.py
- inference.py
- model.py
- requirements.txt
- train.py

Step 6—Make the following changes in the config.py according to our custom dataset of Facemask:

- (a) Ensure the training and validation paths are correctly set (red rectangle)

```
config.py x
config.py > ...
1 import torch
2
3 BATCH_SIZE = 8 # Increase / decrease according to GPU memeory.
4 RESIZE_TO = 300 # Resize the image for training and transforms.
5 NUM_EPOCHS = 50 # Number of epochs to train for.
6 NUM_WORKERS = 4 # Number of parallel workers for data loading.
7
8 DEVICE = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
9
10 # Training images and XML files directory.
11 TRAIN_DIR = 'C:\\Users\\ForAI\\OneDrive\\Desktop\\DL task\\my_ssd_implementation\\data\\Train\\Train\\JPEGImages'
12
13 # Validation images and XML files directory.
14 VALID_DIR = 'C:\\Users\\ForAI\\OneDrive\\Desktop\\DL task\\my_ssd_implementation\\data\\Val\\Val\\JPEGImages'
15
16 # # Classes: 0 index is reserved for background.
17 CLASSES = ['__background__', 'with_mask', 'without_mask', 'mask_wearred_incorrect']
18
19
20 NUM_CLASSES = len(CLASSES)
21
22 # Whether to visualize images after creating the data loaders.
23 VISUALIZE_TRANSFORMED_IMAGES = False
24
25 # Location to save model and plots.
26 OUT_DIR = 'outputs'
```

- (b) In the CLASSES we will be specifying the correct labels for our custom dataset (blue rectangle)
- (c) Also make sure to change the “BATCH_SIZE” according to your GPU therefore if your GPU is weak then lower the “BATCH_SIZE” to avoid further errors

Step 7(a)— Next we define the model.py using resnet34 backbone:

- (a) Import libraries:

```
import torchvision
import torch.nn as nn

from torchvision.models.detection.ssd import (
    SSD,
    DefaultBoxGenerator,
    SSDHead
)
```

(b) Define the model class:

```
def create_model(num_classes=4, size=300, nms=0.45):
    model_backbone = torchvision.models.resnet34(
        weights=torchvision.models.ResNet34_Weights.DEFAULT
    )
    conv1 = model_backbone.conv1
    bn1 = model_backbone.bn1
    relu = model_backbone.relu
    max_pool = model_backbone.maxpool
    layer1 = model_backbone.layer1
    layer2 = model_backbone.layer2
    layer3 = model_backbone.layer3
    layer4 = model_backbone.layer4
    backbone = nn.Sequential(
        conv1, bn1, relu, max_pool,
        layer1, layer2, layer3, layer4
    )
    out_channels = [512, 512, 512, 512, 512, 512]
    anchor_generator = DefaultBoxGenerator(
        [[2], [2, 3], [2, 3], [2, 3], [2], [2]],
    )
    num_anchors = anchor_generator.num_anchors_per_location()
    head = SSDHead(out_channels, num_anchors, num_classes)
    model = SSD(
        backbone=backbone,
        num_classes=num_classes,
        anchor_generator=anchor_generator,
        size=(size, size),
        head=head,
        nms_thresh=nms
    )
    return model
```

Note: make sure the “num_classes” is correctly configured e.g in our case since we have 4 labels hence “num_classes=4”

(c) Next define the main method:

```
if __name__ == '__main__':
    model = create_model(4, 300)
    print(model)
    # Total parameters and trainable parameters.
    total_params = sum(p.numel() for p in model.parameters())
    print(f"total_params:{total_params} total parameters.")
    total_trainable_params = sum(
        p.numel() for p in model.parameters() if p.requires_grad
    )
    print(f"total_trainable_params:{total_trainable_params} training parameters.")
```

Make sure you pass the correct number of classes to the model (red rectangle)

Step 7(b)— (optional)→define the model.py using VGG16 backbone:

(a) Import Libraries

```
import torchvision
from torchvision.models.detection.ssd import SSDClassificationHead
from torchvision.models.detection import _utils
from torchvision.models.detection import SSD300_VGG16_Weights
```

(b) Define the model class:

```
def create_model(num_classes=4, size=300):
    # Load the Torchvision pretrained model.
    model = torchvision.models.detection.ssd300_vgg16(
        weights=SSD300_VGG16_Weights.COCO_V1
    )
    # Retrieve the list of input channels.
    in_channels = _utils.retrieve_out_channels(model.backbone, (size, size))
    # List containing number of anchors based on aspect ratios.
    num_anchors = model.anchor_generator.num_anchors_per_location()
    # The classification head.
    model.head.classification_head = SSDClassificationHead(
        in_channels=in_channels,
        num_anchors=num_anchors,
        num_classes=num_classes,
    )
    # Image size for transforms.
    model.transform.min_size = (size,)
    model.transform.max_size = size
    return model
```

Note: make sure the “num_classes” is correctly configured e.g in our case since we have 4 labels hence “num_classes=4”

(c) Next define the main method:

```
if __name__ == '__main__':
    model = create_model(4, 640)
    print(model)
    # Total parameters and trainable parameters.
    total_params = sum(p.numel() for p in model.parameters())
    print(f"{total_params:,} total parameters.")
    total_trainable_params = sum(
        p.numel() for p in model.parameters() if p.requires_grad
    )
    print(f"{total_trainable_params:,} training parameters.")
```

Make sure you pass the correct number of classes to the model (red rectangle)

Step 8— Next we will be making a couple of changes in the `datasets.py`:

- (a) Import these additional libraries in the script:

```
import matplotlib.pyplot as plt
```

- (b) Replace the “`visualize_sample`” method with the following code:

```
# function to visualize a single sample
def visualize_sample(image, target):
    # Convert from RGB (which OpenCV uses) to BGR (which matplotlib uses).
    image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    # Plot the image
    plt.imshow(image_rgb)
    plt.axis('off') # Hide axes

    # Add bounding boxes and labels
    for box_num in range(len(target['boxes'])):
        box = target['boxes'][box_num]
        label = CLASSES[target['labels'][box_num]]
        plt.gca().add_patch(
            plt.Rectangle(
                (box[0], box[1]),
                box[2] - box[0],
                box[3] - box[1],
                linewidth=2,
                edgecolor='r',
                facecolor='none'
            )
        )
        plt.text(
            box[0],
            box[1] - 5,
            label,
            color='r',
            fontsize=10,
            weight='bold'
        )


    plt.show() # Display the image with annotations
```

Step 9— Now we will start the training by using the following command in the terminal:

```
\DL task\my_ssd_implementation> python train.py
```

Step 10—Saving the training losses, mAP plots & weights:

- (a) After the training has completed the loss curves, mAP score plots and the best weights will be saved in the following “outputs” folder:



Note: use the “best_model.pth” weights when evaluating the model

Step 11—Next evaluate the model (calculate mAP score) on your test dataset by running the following command in terminal:

```
\my ssd implementation> python eval.py
```

Result after the successful execution of eval.py:

```
(ssd_custom) PS C:\Users\ForAI\OneDrive\Desktop\DL task\my_ssd_implementation> python eval.py
Validating
100%|██████████| 1/1 [00:00<00:00]
mAP_50: 65.761
mAP_50_95: 39.558
```


Step 12—Now we will be making some changes in inference.py:

- (a) If your dataset has 10 classes then you should have 10 color definitions but since we have 4 classes hence we have defined 4 colors as shown below in the inference.py script:

```
inference.py •
inference.py > ...
22 )
23 parser.add_argument(
24     '--imgsz',
25     default=None,
26     type=int,
27     help='image resize shape'
28 )
29 parser.add_argument(
30     '--threshold',
31     default=0.25,
32     type=float,
33     help='detection threshold'
34 )
35 args = vars(parser.parse_args())
36
37 os.makedirs('inference_outputs/images', exist_ok=True)
38
39 #make change here must
40 COLORS = [[0, 0, 0], [255, 0, 0], [0,255,0], [0,0,255]]
41
42 #make change here must
43
44
45 # Load the best model and trained weights.
46 model = create_model(num_classes=NUM_CLASSES, size=300)
47 checkpoint = torch.load('outputs/best_model.pth', map_location=DEVICE)
```

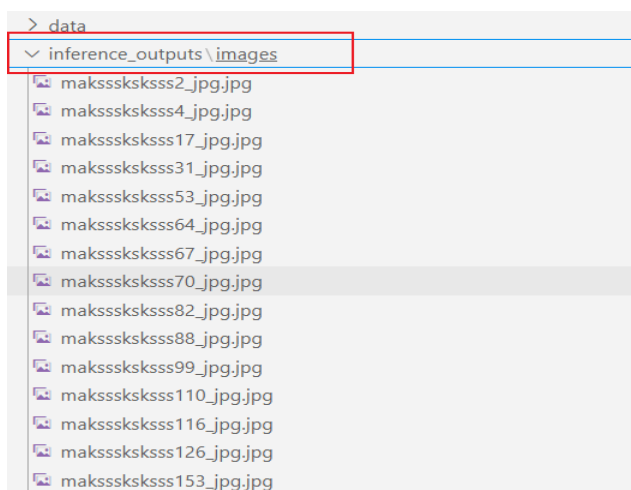
if we have m classes then
we will have m colors as
well as shown here





- (b) Now run the inference.py by using the following command (make sure you correctly specify the test images path):

```
python inference.py --input data/Test/Test/JPEGImages/
```

- (c) After successful execution of inference.py script the predictions on the test set will be saved in the following directory:




Step 13—Visualize your model predictions by viewing some sample images:




Step 14— Evaluating our SSD model with the VGG16 and resnet-34 backbones:

- (a) After 50 epochs we got the following mAP and training loss plots for the resnet-34:



- (a) Similarly, After 50 epochs we got the following mAP and training loss plots for the VGG16:



Conclusion: *VGG16 scores higher mAP (mAP=65.761) than resnet-34 on this dataset for 50 epochs*

Making Improvements to the PyTorch SSD Model with Custom Backbone

There are a few ways to make the model even better.

- *We can start by training it on more data.*
- *We can also try a larger backbone like ResNet50, ResNet101 which may prove to be a better feature extractor.*
- *Also, we can add FPN (Feature Pyramid Network) which can help the backbone a lot when dealing with small objects.*

Final Notes to avoid Potential errors:

- *When using VGG16 model or resnet-34 model make sure you name your model script as model.py and not model_resnet.py or model_vgg16.py*